



codurance.com

GitLab CI Training

Goal

Understand how GitLab Ci pipelines work

What you'll learn

- Set up a CI pipeline
- Options to configure pipelines
- Why you should do this

What is CI?

Continuous Integration

Automating the integration of code changes into a single software project.

What is CI?

Continuous Integration

Automating the integration of code changes into a single software project.



GitLab CI advantages

- Ease of configuration
- Source code security
- Pipeline automation
- GitLab native integration
- Deployment scheduling

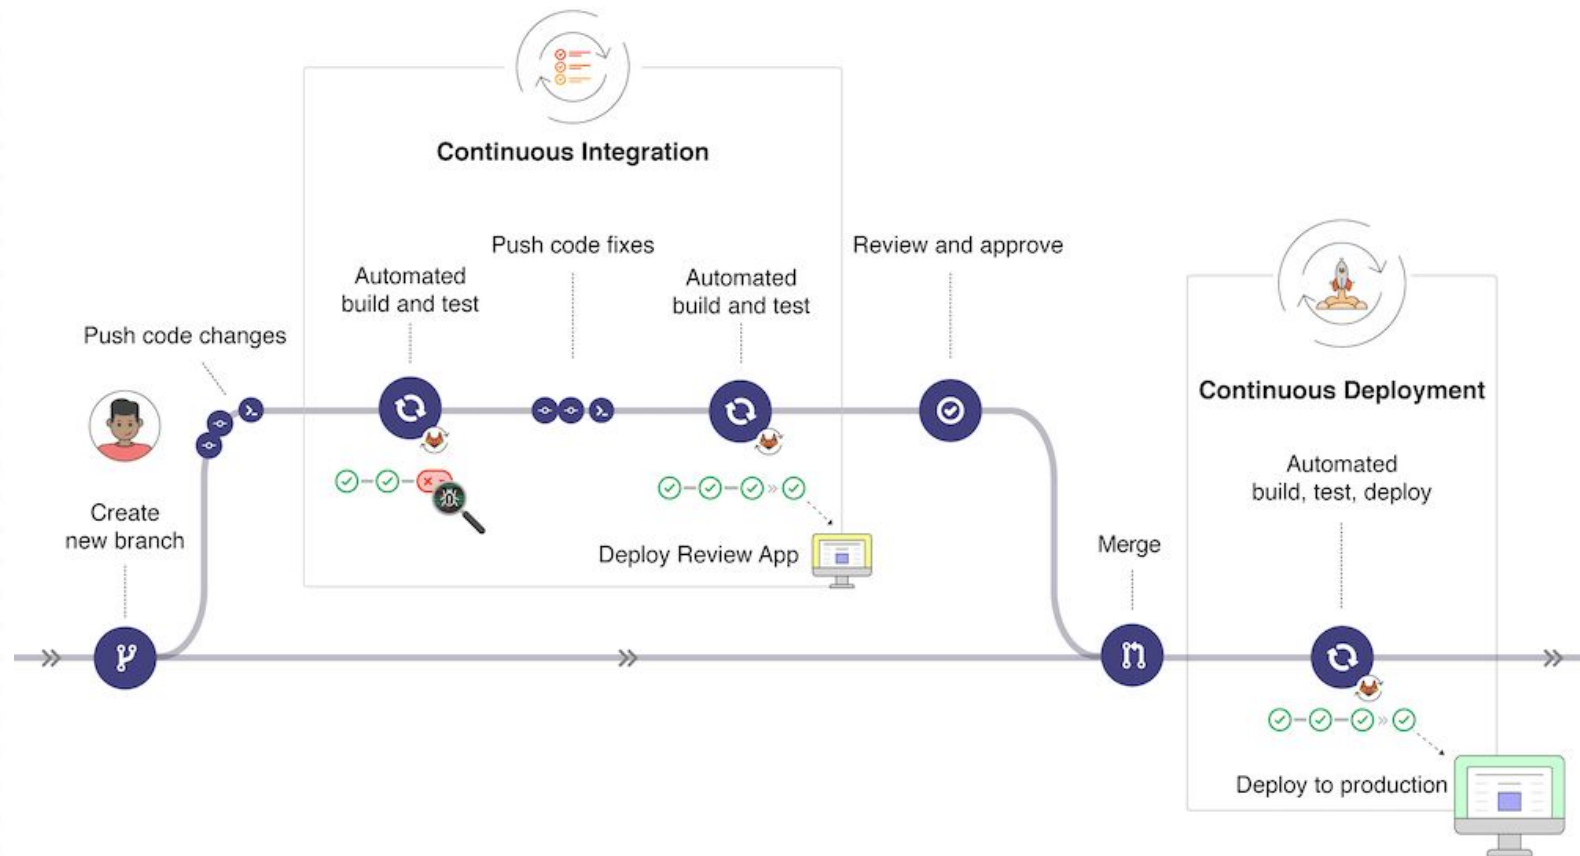
GitLab CI workflow

GitLab CI Workflow

GitLab CI/CD: GitLab continuous methods

- Continuous Integration
- Delivery
- Deployment

GitLab CI Workflow



Pipelines

Pipelines

- A series of automated processes
- Move code updates from VCS to production




Pipelines

Jobs




Stages

Pipelines



Pipelines



Pipelines



Pipelines

1. Build



Pipelines

1. Build
2. Test



Pipelines

1. Build
2. Test
3. Staging




Pipelines

1. Build
2. Test
3. Staging
4. Production



Pipelines

 Menu

G greetings-be

Project information

Repository

Issues 0

Merge requests 0

CI/CD

Pipelines

Editor

Jobs

Schedules

Security & Compliance

Deployments

Packages & Registries

Infrastructure

Monitor

Analytics

Wiki

Snippets


Settings

Platform Engineering > gitlab-ci and argocd workflow > greetings-be > Pipelines

All 1 Finished Branches Tags

Clear runner caches CI lint Run pipeline

Filter pipelines Show Pipeline ID

Status	Pipeline	Triggerer	Stages
<div>passed</div> <div>00:02:34</div> <div>1 week ago</div>	<div>Add .gitlab-ci.yml comments</div> <div>#599543658 main -> b79b90a6</div>		<div>✓✓✓✓</div> <div>▶ ⋮</div>

GitLab CI Training

21


Jobs

Jobs

- Most fundamental element of a pipeline.
- Defined with constraints.
- Top-level elements with an arbitrary name
- Must contain *at least* the [script](#) clause.



Jobs



GitLab Runner

An application that works with GitLab CI/CD to run jobs in a pipeline.



GitLab Runner

Runners delegate the actual execution of jobs to *executors*




GitLab Executor

An executor determines the environment each job runs in.




GitLab Runners + Executors







Executors

Executor	SSH	Shell	VirtualBox	Parallels	Docker	Kubernetes	Custom
Clean build environment for every build	X	X	✓	✓	✓	✓	conditional (4)
Reuse previous clone if it exists	✓	✓	X	X	✓	X	conditional (4)
Runner file system access protected (5)	✓	X	✓	✓	✓	✓	conditional
Migrate runner machine	X	X	partial	partial	✓	✓	✓
Zero-configuration support for concurrent builds	X	X (1)	✓	✓	✓	✓	conditional (4)
Complicated build environments	X	X (2)	✓ (3)	✓ (3)	✓	✓	✓
Debugging build problems	easy	easy	hard	hard	medium	medium	medium




Jobs



Jobs




Jobs



Pipeline Configuration

Pipeline configuration

 Menu

Mauro Chojrin > My first GitLab Project > Pipelines > #605978526

M My first GitLab Project

Project information

Repository

Issues 0

Merge requests 0

CI/CD

Pipelines

Editor

Jobs

Schedules

Security & Compliance

Deployments

Packages & Registries

Infrastructure

Monitor

Analytics

Wiki

Snippets

Settings

passed

 Pipeline #605978526 triggered 4 days ago by Mauro Chojrin

Delete

Adding a test stage

3 jobs for main in 17 seconds

latest


793ff524

No related merge requests found.

Pipeline Needs **Jobs 3** Tests 0

Status	Job	Stage	Name	Duration	Coverage
passed	#2820683786 main ~ 793ff524	test	test-job	00:00:08 4 days ago	
passed	#2820680393 main ~ 793ff524	salute	build-job	00:00:08 4 days ago	
Failed	#2820680392 main ~ 793ff524	test	test-job	00:00:09 4 days ago	

Pipeline configuration

 Menu

Search GitLab

M

My first GitLab Project

Project information

Repository

Issues

Merge requests

CI/CD

Pipelines

Editor

Jobs

Schedules

Security & Compliance

Deployments

Packages & Registries

Infrastructure

Monitor

Analytics

Wiki

Snippets

Settings

Mauro Chojrin > My first GitLab Project > Jobs > #2820683786

passed

Job test-job triggered 4 days ago by Mauro Chojrin

1 Running with gitlab-runner 15.2.1 (32fc1585)

2 on mauro-codurance His2Aq1Y

3 **Preparing the "docker" executor** 00:03

4 Using Docker executor with image alpine:latest ...

5 Pulling docker image alpine:latest ...

6 Using docker image sha256:d7d3d98c851ff3a95dbcb70ce09d186c9aaf7e25d48d55c0f99aae360aecfd53 for alpine:latest with digest alpine@sha256:7580ece7963bfa863801466c0a488f11c86f85d9988051a9f9c68cb27f6b7872 ...

7 **Preparing environment** 00:01

8 Running on runner-his2aqly-project-38373599-concurrent-0 via mauro-codurance...

9 **Getting source from Git repository** 00:02

10 Fetching changes with git depth set to 20...

11 Reinitialized existing Git repository in /home/gitlab-runner/build/mauro.chojrin/my-first-gitlab-project/.git/

12 **Checking out 793ff524 as main...**

13 Skipping Git submodules setup

14 **Executing "step_script" stage of the job script** 00:00

15 Using docker image sha256:d7d3d98c851ff3a95dbcb70ce09d186c9aaf7e25d48d55c0f99aae360aecfd53 for alpine:latest with digest alpine@sha256:7580ece7963bfa863801466c0a488f11c86f85d9988051a9f9c68cb27f6b7872 ...

16 **\$ echo "Performing really complex tests!"**

17 Performing really complex tests!

18 **Cleaning up project directory and file based variables** 00:00

19 **Job succeeded**

test-job

Duration: 8 seconds

Finished: 4 days ago

Queued: 0 seconds

Timeout: 1h (from project)

Runner: #16817953 (His2Aq1Y) mauro-codurance

Commit 793ff524

Adding a test stage

Pipeline #605978526 for main

test

test-job

test-job

GitLab CI/CD

1. Application code hosted in a Git repository.
2. A file called `.gitlab-ci.yml` in the root of your project

.gitlab-ci.yml

- Define, declaratively, how we want to run our pipeline
 - Scripts you want to run.
 - Other configuration files and templates you want to include.
 - Dependencies and caches.
 - Commands you want to run in sequence and those you want to run in parallel.
 - Location to deploy your application to.
 - Whether to run the scripts automatically or trigger any of them manually.

Exercise: My first pipeline

1. Create a pipeline that returns “Hello, Pipeline”
2. Navigate through the pipelines interface to see the results.

```
1  stages:           # List of stages for jobs, and their order of execution
2  |   - salute
3
4  build-job:        # This job runs in the build stage, which runs first.
5  |   stage: salute
6  |   script:
7  |     - echo "Hello, Pipeline"
```

Optional: Installing a local GitLab Runner in Ubuntu

1. Validate that you have docker installed and can be run in non-root mode
2. Download binary file from <https://gitlab-runner-downloads.s3.amazonaws.com/latest/index.html>
3. Get configuration details from https://gitlab.com/<YOUR_PROJECT>/-/settings/ci_cd
4. Start it up: `sudo gitlab-runner start`
5. Register it using `sudo gitlab-runner register`
6. When prompted, enter the URL and Token provided at step 3
7. Select `docker` as executor
8. default image: `alpine:latest`
9. Reload GitLab's page to check everything is ok.
10. Disable shared runners
11. Clone the repository
12. Add a new file
13. Commit
14. Push
15. Watch the pipeline run

Pipeline visualization

- Pipeline view
- Pipeline ID
- Branch
- Commit
- Triggerer
- Retry job
- Job logs
- Job view
- Job ID

Status	Pipeline	Triggerer	Stages
<div>passed</div> <div>00:00:28</div> <div>6 minutes ago</div>	Update .gitlab-ci.yml file #11734 main -> e516e8d7 latest		<div>passed</div>

All 1 Pending 0 Running 0 Finished 1						
Status	Name	Job	Pipeline	Stage	Duration	Coverage
<div>passed</div>	build-job	#40391 main -> e516e8d7	#11734 by	salute	<div>00:00:28</div> <div>6 minutes ago</div>	

GitLab CI keywords

GitLab CI keywords

Keywords define **how** we want our pipelines and jobs to run.

GitLab CI keywords



Global keywords

default

Set global defaults for other keywords in all the pipeline

```
1  default:
2    image: alpine:latest
3
4  job-1:
5    script:
6      - echo "this is the first job and it's going to run on alpine image"
7
8  job-2:
9    image: maven:3-jdk-17
10   script:
11     - echo "this is the second job and it's going to run on maven image"
12
```

Global keywords

default

Set global defaults for other keywords in all the pipeline

```
1  default:
2    image: alpine:latest
3
4  job-1:
5    script:
6      - echo "this is the first job and it's going to run on alpine image"
7
8  job-2:
9    image: maven:3-jdk-17
10   script:
11     - echo "this is the second job and it's going to run on maven image"
12
```

Keywords:

- after_script
- artifacts
- before_script
- cache
- image
- interruptible
- retry
- services
- tags
- timeout

Global keywords

stages

Stages that contain groups of jobs

```
1  stages:
2    - build
3    - test
4    - deploy
```

Global keywords

stages

- .pre
- build
- test
- deploy
- .post

Job keywords

stage

- Which stage a job runs in
- Defaults to *test*
- Can have more than one value

Job keywords

image

Docker image the job runs in.

```
1  job-1:
2    image: alpine:latest
3    script:
4      - echo "this is a job and it's going to run on an alpine docker image"
```

Job keywords

variables

Custom values passed to jobs.

```
1 job-1:
2   image: alpine:latest
3   variables:
4     DOCKER_IMAGE: alpine
5   script:
6     - echo "this is a job and it's going to run on an $DOCKER_IMAGE docker image"
7
```

Job keywords

script

Commands for the runner to execute.

```
1 job-1:
2   image: alpine:latest
3   variables:
4     DOCKER_IMAGE: alpine
5   script:
6     - echo "this is a job and it's going to run on an $DOCKER_IMAGE docker image"
7     - echo "we can add several lines to the script keyword"
```

Job keywords

when

- Conditions for the job to run
- Boolean values

Job keywords

when

- Conditions for the job to run
- Boolean values
 - a. **on_success**

Job keywords

when

- Conditions for the job to run
- Boolean values
 - a. **on_success**
 - b. **manual**

Job keywords

when

- Conditions for the job to run
- Boolean values
 - a. **on_success**
 - b. **manual**
 - c. **always**

Job keywords

when

- Conditions for the job to run
- Boolean values
 - a. **on_success**
 - b. **manual**
 - c. **always**
 - d. **on_failure**

Job keywords

when

- Conditions for the job to run
- Boolean values
 - a. **on_success**
 - b. **manual**
 - c. **always**
 - d. **on_failure**
 - e. **delayed**

Job keywords

when

- Conditions for the job to run
- Boolean values
 - a. **on_success**
 - b. **manual**
 - c. **always**
 - d. **on_failure**
 - e. **delayed**
 - f. **never**

Job keywords

rules

Select jobs to run based on conditions.

Job keywords

rules

- *if*
- *changes*
- *exists*
- *allow_failure*
- *variables*
- *when*

Job keywords

rules

```
variables:  
  VARIABLE: "manual"
```

```
build-job-conditional:  
  stage: salute  
  script:  
    - echo "This is a conditional manual step"  
  rules:  
    - if: $VARIABLE == "manual"  
      when: manual
```

Job keywords

allow_failure

Continue running the pipeline if that job fails

```
build-job:  
  stage: salute  
  allow_failure: true  
  script:  
    - exit 1
```

Salute



build-job



build-job - failed - (script failure)
(allowed to fail)

Job keywords

allow_failure

Default values:

- `true` for manual jobs.
- `false` for jobs that use `when: manual` inside rules
- `false` in all other cases.

```
build-job:  
  stage: salute  
  allow_failure: true  
  script:  
  - exit 1
```

Salute



build-job



build-job - failed - (script failure)
(allowed to fail)

Job keywords

artifacts

Files and/or directories to save once the job finishes

```
1  job-1:
2    script:
3      - echo "here we build our java package"
4    artifacts:
5      expire_in: 1 hour
6      paths:
7        - target/*.jar # Here we save the jar as an artifact
```

GitLab CI variables

GitLab CI variables

CI/CD variables are a type of environment variable. We can use them to:

- Control the behavior of jobs and pipelines.
- Store values we want to reuse.
- Avoid hard-coding values in our `.gitlab-ci.yml` file.

We can use a set of already predefined variables or we can define our own.

GitLab CI variables

A type of environment variable.

GitLab CI variables

- **Control the behavior** of jobs and pipelines.
- **Store values** we want to reuse.
- **Avoid hard-coding** values in our `.gitlab-ci.yml` file.

Predefined variables

GitLab CI/CD has a default set of predefined CI/CD variables.

Available in every GitLab CI/CD pipeline.

Custom variables: .gitlab-ci.yml

- Defined with *variables* keyword.
- Visible in the code and the repository.

```
1  variables:
2  |   PLATFORM: GitLab # All jobs defined can use it
3  job-1:
4  |   variables:
5  |     COMPANY: "Codurance" # Only the job can use it
6  |   script:
7  |     - echo "I'm taking a $PLATFORM training from $COMPANY"
```

Variable scope

- Job level
 - Only available at the current job
- Pipeline level
 - Available for all jobs in the current pipeline
- Project level
 - Available for all pipelines in the project
 - Possibility to scope by environment
- Group level
 - Available for all pipelines in the projects inside the group
 - Possibility to scope by environment
- Predefined variables
 - Available for all pipelines that run in GitLab

Exercise: Ci Pipeline

Automate the CI of our application, [greetings](#)

Exercise: Ci Pipeline

- Clear the steps to build it,
- Have written tests following best practices.

Exercise: Ci Pipeline

- Reduce the time our developers spend with the CI
- Reduce the human errors due to repetitive manual tasks

Exercise: Ci Pipeline

We are going to use our gained knowledge in GitLab CI.


Exercise: Ci Pipeline

The first step will be to define the pipeline:


- How many stages?
- How many jobs?
- Which ones?
- What are the dependencies between them?

Exercise: Ci Pipeline

In this case we are going to focus on a basic CI pipeline:



Exercise: Ci Pipeline



- Build job to compile and build our package
- Test job to run our tests
- Build image job to build a docker image based on our compiled code
- Deploy job to deploy to our environment

Exercise: Ci Pipeline. Build package

- What is our programming language?
- What docker image should we use as base?
- What artifacts will we need to use later?
- Do we need to store any kind of variables?

Exercise: Ci Pipeline. Test

When creating this job we need to ask ourselves:

- We will again ask all the questions from the build phase, as they will probably align with this one.
- How do I report my tests to GitLab?
- Do I want to have just a separate test job?
 - If you are clear on how to separate your tests (Unit tests, Integration tests...) and execute them separately then go for it! It will help you diagnose later. You can run them all in the same test stage but with different names.
 - If you are not clear on it, as an initial advice the best option is to go for a generic test job and run all your tests. You can always separate them later!

Demo CI pipeline

Advanced elements


Global keywords

include

Include external yaml files in our CI/CD configuration.

```
1  include: '.build-template.yml'  
2
```

Pipeline notifications



Pipeline notifications



Pipeline notifications

The screenshot shows the GitLab web interface for configuring pipeline status email notifications. The left sidebar contains a navigation menu with options like Project information, Repository, Issues, Merge requests, CI/CD, Security & Compliance, Deployments, Packages & Registries, Infrastructure, Monitor, Analytics, Wiki, Snippets, and Settings. The 'Settings' section is expanded, showing 'Integrations' as the active tab. The main content area is titled 'Pipeline status emails' and includes a search bar, a breadcrumb trail, and several configuration options. The 'Enable Integration' section has a checked 'Active' checkbox. The 'Recipients' section has a text input field for email addresses. The 'Notify only broken pipelines' section has a checked checkbox. The 'Branches for which notifications are to be sent' section has a dropdown menu set to 'Default branch'. At the bottom, there are three buttons: 'Save changes', 'Test settings', and 'Cancel'. A mouse cursor is visible over the 'Test settings' button.

Menu

g greetings-be

- Project information
- Repository
- Issues 0
- Merge requests 0
- CI/CD
- Security & Compliance
- Deployments
- Packages & Registries
- Infrastructure
- Monitor
- Analytics
- Wiki
- Snippets
- Settings
 - General
 - Integrations**
 - Webhooks
 - Repository
 - CI/CD
 - Packages & Registries
 - Pages
 - Monitor
 - Usage Quotas

Platform Engineering > gitlab-ci and argo workflow > greetings-be > Integration Settings > Pipeline status emails

Search page

Pipeline status emails

Enable Integration

☒ Active

Recipients

Comma-separated list of email addresses.

Notify only broken pipelines

☒ Notify only broken pipelines

Branches for which notifications are to be sent

Default branch

Save changes Test settings Cancel

« Collapse sidebar

Job keywords

environment

The environment a job runs at.

Custom variables: Project/Group settings

We can add CI/CD variables to a group/project's settings.

Only project members with the Maintainer role can add or update project CI/CD variables. To keep a CI/CD variable secret, we will put it in the project settings, not in the `.gitlab-ci.yml` file.

We can add simple variables or full files as variables, such as SSH keys or parameters files.

Custom variables: Project/Group settings

To add or update variables in the project settings:

- Go to your project's Settings > CI/CD > Variables
- Select the Add Variable button and fill in the details:
 - Key: the key for our variable
 - Value: the content of our variable
 - Type: File or Variable.
 - Environment scope: Optional. All, or specific environments.
 - Protect variable Optional. If selected, the variable is only available in pipelines that run on protected branches or protected tags.
 - Mask variable Optional. If selected, the variable's Value is masked in job logs. The variable fails to save if the value does not meet the masking requirements.

Variable precedence

1. These all have the same (highest) precedence:
 - Trigger variables.
 - Scheduled pipeline variables.
 - Manual pipeline run variables.
 - Variables added when creating a pipeline with the API.
2. Project variables.
3. Group variables.
4. Instance variables.
5. Inherited variables.
6. Variables defined in jobs in the `.gitlab-ci.yml` file.
7. Variables defined outside of jobs (globally) in the `.gitlab-ci.yml` file.
8. Deployment variables.
9. Predefined variables.

Pipeline images in our own GitLab Container Registry

M

Job keywords

needs

This keyword will allow us to execute jobs out of order, ignoring stage ordering and allowing some jobs to run without waiting all the jobs from a previous stage to complete.

If we leave this keyword empty it will mean that the job is set to start as soon as the pipeline is created.

Job keywords

dependencies

We will use the dependencies keyword to define a list of jobs to fetch artifacts from. We can also set a job to download no artifacts at all.

If we do not use dependencies, all artifacts from previous stages are passed to each job.

If we leave the keyword empty we will be configuring the job to not download any artifacts.

Job keywords

before_script / *after_script*

We will use these keywords to define an array of commands that should run before/after each job's *script* commands.

- *before*: they are concatenated with the *script* commands and they run in a single shell
- *after*: they run even if the job fails. These scripts execute in a new shell, separate from the *before_script* and *script* commands. That means they don't have access to the changes done in those commands.

Trigger keyword

M

Merge requests and working with features

M

Code coverage

—

M



Job keywords

coverage

This keyword will be used

Badges

M



Security configuration

M

GitLab CI registries

GitLab Container Registry

The GitLab Container Registry is a secure and private registry for container images. It's built on open source software and completely integrated within GitLab.

The stored Images follow this naming convention:

<registry URL>/<namespace>/<project>/<image>

We will use GitLab CI/CD to create and publish images from the pipelines. In order to do that we will have to authenticate with the Container Registry. That's when the predefined variables enter in place with these two options:

- `docker login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD $CI_REGISTRY`
- `docker login -u $CI_REGISTRY_USER -p $CI_JOB_TOKEN $CI_REGISTRY`

Pull images from GitLab Container Registry

1. Create access token

We will access into our GitLab profile and we will create an access token with permissions to read and write into the registry:

2. Log into the docker registry

We will login into the docker registry with our GitLab username and our newly created token:

```
amorell@LAPTOP-BIT56HGB:~$ docker login -u alex [MYREGISTRY]
```


Password:

Login Succeeded

3. Pull docker image

We will pull the docker image with the standard commands:

```
docker pull [MYREGISTRY]/[MYIMAGE]:[MYTAG]
```



GitLab Package Registry

With the GitLab Package Registry, we can use GitLab as a private or public registry for a variety of supported package managers. We can also publish and share packages, which can be consumed as a dependency in downstream projects.

In order to view packages for your project or group:

1. Go to the project or group.
2. Go to Packages & Registries > Package Registry.

We will be able to search, sort, and filter packages on this page, and also share your search results by copying and pasting the URL from our browser.

The package manager supports several formats:

Package type	GitLab version	Status
Maven	11.3+	GA
npm	11.7+	GA
NuGet	12.8+	GA
PyPI	12.10+	GA
Generic packages	13.5+	GA
Composer	13.2+	Beta
Conan	12.6+	Beta
Helm	14.1+	Beta
Debian	14.2+	Alpha
Go	13.1+	Alpha
Ruby gems	13.10+	Alpha

Exercise: Advanced Ci Pipeline. Build image

When creating this job we need to ask ourselves:

- Do I need any Docker specific variables?
- When do I want to build the image, on main branch after merge requests or also on features?
- Do I have a Dockerfile already, and where is it stored?
- Where do I want to store my built Image? Will it be my own registry or GitLab registry?

Resources

Additional resources

<https://docs.gitlab.com/ee/ci/>

<https://docs.gitlab.com/ee/ci/pipelines/>

<https://docs.gitlab.com/ee/ci/yaml/>

https://docs.gitlab.com/ee/ci/variables/predefined_variables.html

<https://nexocode.com/blog/posts/understanding-principles-of-gitlab-ci-cd-pipelines/>

<https://github.com/firecow/gitlab-ci-local>



Thank you



Fran Iglesias
Software Craftperson

✉ fran.iglesias@codurance.com

in [Fran Iglesias](#)



Mauro Chojrin
Software Craftperson

✉ mauro.chojrin@codurance.com

in [Mauro Chojrin](#)